
Active Heirarchical Metric Learning

Nicolas Beltran

Department of Computer Science
Columbia University
New York City, NY 10027
nb2838@columbia.edu

Ketan Jog

Department of Computer Science
Columbia University
New York City, NY 10027
kj2473@columbia.edu

Abstract

Many problems require a well defined notion of a distance between points in space. Constructing or finding such a measure falls into the field of metric learning. Although many algorithms exist in the field when a learner has access to a fixed dataset, there is room for improvement in terms of samples efficiency that the learner needs to know, imposition of desired structure, especially when the data appears in an *online* manner. We propose a project that reduces the problem of online/active metric learning to bandits. In case our plan turn out to be too ambitious, we have a fallback - an empirical investigation of some algorithms that have dealt with the problem in an online setting or in situations where the learner can selective query the points that it wants to know information about.

1 Introduction

2 Related Work

3 Long-term goals

4 Preliminaries

5 Problem Statement

We consider two different problems. The first problem consists of making a series of sequential predictions while learning a similarity measure. We refer to this problem as online similarity prediction. The second problem consists of learning a similarity measure while querying points in space. We refer to this problem as active similarity learning.

5.1 Online similarity learning

We consider an online similarity learning problem played over T rounds. At round t the environment samples K pairs of points $(x_{t,k}, y_{t,k}) \in \mathbb{R}^{2n}$. The agent then chooses pair $k_t \in [K]$ and is given a reward $r_{t,k_t} \in \{1, -1\}$. We assume that there exists some similarity function unknown to the agent $\phi : \mathbb{R}^{2n} \rightarrow \{-1, 1\}$ and that the rewards are such that if at time $t \in [T]$ the agent chooses pair $(x_{t,k}, y_{t,k})$ then the reward is $\phi(x_{t,k}, y_{t,k})$.

As usual we define the regret as

$$R_T = \mathbb{E} \left[\sum_{t=1}^T \phi(x_{t,k_t^*}, y_{t,k_t^*}) - \phi(x_{t,k_t}, y_{t,k_t}) \right]$$

where $k_t^* = \operatorname{argmax}_{k \in [K]} \phi(x_{t,k}, y_{t,k})$

5.2 Active similarity learning

We assume that the learner has access to a dataset $D = \{x_i \in \mathbb{R}^n | i \in [N]\}$ of unlabeled points and that there exists some function $\phi : \mathbb{R}^{2n} \rightarrow \{-1, 1\}$ which the learner is trying to learn. The learner

can query T pairs of points in this set D to obtain a dataset $D_T = \{(x_t, y_t, r_t) \mid t \in [T]\}$ where x_t and y_t are the points in space and r_t is their similarity. The learner maintains and estimate $\hat{\phi}_t \in \mathcal{F}$ of ϕ , where \mathcal{F} is its function class and we denote the loss between an estimate $\hat{\phi}$ and ϕ as

$$\mathcal{L}(\phi, \hat{\phi}) = \mathbb{E}_{(x,y) \sim \mathcal{D} \times \mathcal{D}} [(\hat{\phi}(x, y) - \phi(x, y))^2]$$

In this problem the goal is to find $\min_{\phi \in \mathcal{F}} \mathcal{L}(\hat{\phi}_T, \phi)$.

6 Description of the algorithms

We provide 4 algorithms which can be grouped into neural and linear bandit version. We refer to our algorithms as OnSim-LinUCB, ActSim-LinUCB, OnSim-NeuralUCB, ActSim-NeuralUCB, depending on whether the bandit algorithm they are based on and the problem they are trying to solve. The algorithms are all very similar but we provide them in full for completeness.

6.1 Online similarity learning

6.1.1 OnSim-LinUCB

By assuming a linear structure in the similarity function, our first algorithm performs a straightforward reduction of the online learning problem to that of regular contextual linear bandits as described in [1]. Mathematically, we model the similarity of two points as $\phi(x, y) = x^\top A y$. This is reasonable if we consider that

$$\phi(x, y) = x^\top A y = \sum_{i=1}^n \sum_{j=1}^n x_i y_j A_{i,j}$$

which is linear in A and thus allows us to use the framework of LinUCB for our problem. Attending to this fact, one can see that Algorithm 1 is almost identical to the first algorithms in [1].

Algorithm 1: OnSim-LinUCB

Input: Rounds T and exploration parameter α

- 1 $A \leftarrow I_{n^2}$;
- 2 $b \leftarrow 0_{n^2}$;
- 3 **for** $t \in [T]$ **do**
- 4 $\theta_t \leftarrow A^{-1}b$ Observe K pairs of vectors $x^k \in \mathbb{R}^n, y^k \in \mathbb{R}^n$;
- 5 Create $z_{t,k} = (x_1^k, y_1^k, x_1^k, y_1^k, \dots, x_n^k y_{n-1}^k, x_n^k, y_n^k)$;
- 6 **for** $k \in [K]$ **do**
- 7 $p_{t,a} \leftarrow \theta_t^\top z_{t,k} + \alpha \sqrt{z_{t,k}^\top A^{-1} z_{t,k}}$
- 8 **end**
- 9 Choose action $k_t = \operatorname{argmax}_a p_{t,a}$ with ties broken arbitrarily;
- 10 Observe payoff $r_t \in \{-1, 1\}$;
- 11 $A \leftarrow A + z_{t,k_t} z_{t,k_t}^\top$;
- 12 $b \leftarrow z_{t,k_t} r_t$;
- 13 **end**

6.1.2 OnSim-NeuralUCB

The expressive power of the above framework is quite limited because of the fact that Algorithm 1 assumes that similarity is the result of a dot product in which both points are independently mapped to a new representation via a linear transformation. This will obviously not be true for many datasets. To overcome these limitations we adapt NeuralUCB [4] to our circumstances.

Formally, we assume that

$$\phi(x, y) = \cos(f(x; \theta), f(y; \theta)) = \frac{\langle f(x; \theta), f(y; \theta) \rangle}{\|f(x; \theta)\|_2 \|f(y; \theta)\|_2}$$

¹ Here $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ is a simple neural network of the form

$$f(x; \theta) = b_L + W_L \sigma(b_{L-1} + W_{L-1} \sigma(\dots \sigma(b_1 + W_1 x)))$$

for some $L \in \mathbb{N}_{\geq 2}$, $W_1 \in \mathbb{R}^{d \times n}$, $W_L \in \mathbb{R}^{m \times d}$, $W_l \in \mathbb{R}^{d \times d}$ if $l \in [2, L-1]$, $b_L \in \mathbb{R}^m$, $b_l \in \mathbb{R}^d$ for $l \neq L$, and where $\sigma(x) = \max\{0, x\}$ applied component wise. We use θ to denote the flattened vector of $(W_L, b_L, \dots, W_1, b_1)$. Using this notation OnSim-NeuralUCB is described below in Algorithm 2. We use p to denote the total number of trainable parameters in the neural network.

Algorithm 2: OnSim-NeuralUCB

Input : Rounds T , exploration parameter α , τ_r frequency of resets, τ_T frequency of training, E epochs for training, b_s batch size for training, ϵ learning rate.

```

1  $A \leftarrow I_p$ ;
2 for  $t \in [T]$  do
3   Observe  $K$  pairs of vectors  $x_t^k \in \mathbb{R}^n$ ,  $y_t^k \in \mathbb{R}^\kappa$ ;
4   for  $k \in [K]$  do
5      $p_{t,k} \leftarrow \phi(x, y) + \alpha \sqrt{(\nabla_\theta \phi(x_t^k, y_t^k))^\top A^{-1} \nabla_\theta \phi(x_t^k, y_t^k)}$ ;
6   end
7   Choose pair  $k_t = \operatorname{argmax}_k p_{t,k}$  with ties broken arbitrarily;
8   Observe payoff  $r_t \in \{-1, 1\}$ ;
9   if  $t \bmod \tau_r = 0$  then
10     $\theta \leftarrow \text{Train}(\epsilon, E, \{(x_i^{k_i}, y_i^{k_i}, r_i)\}_{i=1}^t, b_s)$ ;
11  end
12  if  $t \bmod \tau_T = 0$  then
13     $A \leftarrow I_p$ 
14  end
15   $A \leftarrow A + \nabla_\theta \phi(x_t^{k_t}, y_t^{k_t})(\nabla_\theta \phi(x_t^{k_t}, y_t^{k_t}))^\top$ ;
16 end
```

Intuitively, this algorithm acts optimally with the current set of parameters with a bonus for exploration, and every so often it stops back to train the network. "Train" can be found in the appendix as 5 but it simply describes the normal process of training a neural network with MSE loss.

Our algorithm is very similar to that proposed in [4] but it differs in the following aspects: 1. We adopt a siamese neural network architecture unlike the paper. 2. We add bias 3. We use cosine similarity function in our last layer 4. We restart the the matrix A after every $t \bmod \tau_T = 0$ iterations 5. We use a constant exploration parameter. We propose these changes because we believe they are more sensible for our particular application (and do provide better empirical performance) but they destroy any theoretical guarantees. In [4] these assumptions alongside arguments related to the neural tangent kernel matrix [2] are used to provide a $\tilde{O}(\sqrt{T})$, bound on the regret of algorithm.

6.2 Active Similarity Learning

Despite the differences with online similarity learning, for the active version of the problem we modify the above algorithms only slightly. The idea is to use the optimism bonus as a measure of uncertainty and ignore completely the reward.

Using this framework, each time we have to make a query we can sample a subset of points in the unlabeled dataset D and then we reveal the label of the pair with the highest uncertainty as measured by the bonus. We emphasize that r_t now references a label rather than a reward. Algorithms 3 and 4 describe these processes formally. In Algorithm 3 the most important change is line 5 and in algorithm 4 the most important change is line 7.

¹It is also possible to assume that we don't normalize the last inner product, we refer to this version and unnormalized NeuralUCB. We provide some comparisons of this version also below.

Algorithm 3: Active-NeuralUCB

Input : Queries T , exploration parameter α , τ_r frequency of resets, τ_T frequency of training, E epochs for training, b_s batch size for training, ϵ learning rate.

```
1  $A \leftarrow I_p$ ;
2 for  $t \in [T]$  do
3   Sample  $K$  pairs of vectors  $x_t^k \in \mathbb{R}^n, y_t^k \in \mathbb{R}^n$  from dataset  $D$ ;
4   for  $k \in [K]$  do
5      $p_{t,k} \leftarrow \alpha \sqrt{(\nabla_{\theta} \phi(x_t^k, y_t^k))^{\top} A^{-1} \nabla_{\theta} \phi(x_t^k, y_t^k)}$ ;
6   end
7   Choose pair  $k_t = \operatorname{argmax}_k p_{t,k}$  with ties broken arbitrarily;
8   Query label  $r_t \in \{-1, 1\}$ ;
9   if  $t \bmod \tau_r = 0$  then
10     $\theta \leftarrow \text{Train}(\epsilon, E, \{(x_i^{k_i}, y_i^{k_i}, r_i)\}_{i=1}^t, b_s)$ ;
11  end
12  if  $t \bmod \tau_T = 0$  then
13     $A \leftarrow I_p$ 
14  end
15   $A \leftarrow A + \nabla_{\theta} \phi(x_t^{k_t}, y_t^{k_t})(\nabla_{\theta} \phi(x_t^{k_t}, y_t^{k_t}))^{\top}$ ;
16 end
```

7 Empirical Evaluation

Below we describe the set of experiments that we performed. We test separately NeuralUCB and LinUCB based algorithms as we logically expect them to perform very differently due to the limitations of a linear approach.

Algorithm 4: Active-LinUCB

Input : Rounds T and exploration parameter α

```
1  $A \leftarrow I_{n^2}$ ;
2  $b \leftarrow 0_{n^2}$ ;
3 for  $t \in [T]$  do
4    $\theta_t \leftarrow A^{-1}b$  Observe  $K$  pairs of vectors  $x^k \in \mathbb{R}^n, y^k \in \mathbb{R}^n$ ;
5   Create  $z_{t,k} = (x_1^k y_1^k, x_1^k y_1^k, \dots, x_n^k y_{n-1}^k, x_n^k y_n^k)$ ;
6   for  $k \in [K]$  do
7      $p_{t,k} \leftarrow \alpha \sqrt{z_{t,k} A^{-1} z_{t,k}^{\top}}$ 
8   end
9   Choose action pair  $k_t = \operatorname{argmax}_k p_{t,k}$  with ties broken arbitrarily.;
10  Observe label  $r_t \in \{-1, 1\}$ ;
11   $A \leftarrow A + z_{t,k_t} z_{t,k_t}^{\top}$ ;
12   $b \leftarrow z_{t,k_t} r_t$ ;
13 end
```

7.1 NeuralUCB

Dataset Although we tested on various datasets including MNIST, CIFAR-10 and crescent moons, we found that the algorithm was impractical to use on the first two due to the expensive computation of the matrix operation. To reduce the dimension we had to apply PCA. However because this process implicitly leads to a better representation useful for our scenarios, we opted not to include the results here an only focus on the performance of crescent moons. However, these experiments can be found on the associated github repo.

Neural Network For our experiments we used a 3 layer neural network with 25 hidden units in each layer and an output dimension of 2. In addition to the description of the algorithm provided above, we also added a dropout layer after each regular layer of the neural network with $p = 0.3$.

Optimizer We attempted to use SGD, Adam, and SGD with momentum with various different hyper parameters. Generally, We found that Adam was the best performing optimizer followed by SGD with momentum. We used a learning rate of 0.001, a momentum of 0.9, and the default hyperparameters in [3].

Additional Parameters We sampled ten datapoints at time for both the online and active version of the problem. Neural networks were trained after 100 queries (or timesteps) and trained for two epochs. Every 4 epochs we reset the matrices to their original states. We rotate the the data the algorithm sees every two time steps. The values reported for losses and accuracy are computed on a heldout set of the data of 1000 datapoints.

7.1.1 Experiment 1: Effectiveness of a linear classifier on learned features

To determine the quality of our embeddings we used an SVM to classify the learned features and observe these values per iteration. We report the results both for SGD with momentum and Adam. Additionally, we also report the performance of NeuralUCB using the unnormalized version of cosine similarity. All graphs contain the mean and a 95% confidence interval computed over 10 runs of the algorithm. The results are shown in Figures (1) and (2)

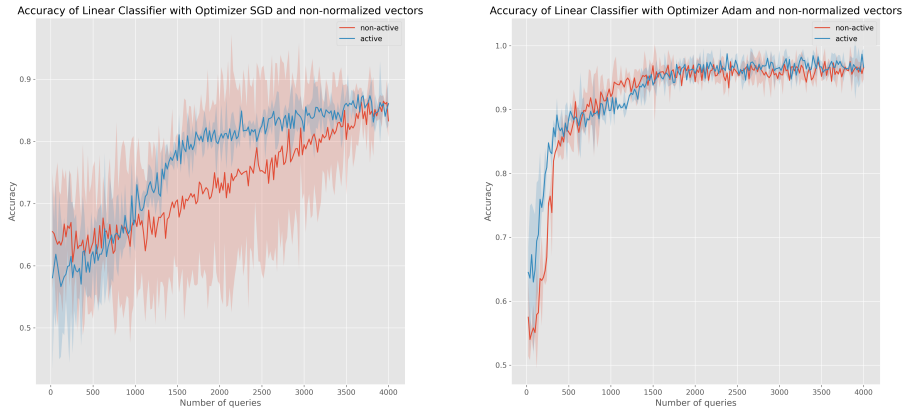


Figure 1: Performance SVM on non-normalized learned features Adam vs SGD




Figure 2: Performance SVM on normalized learned features Adam vs SGD

7.2 Experiment 2: Comparison of L2 loss through time

Using the exact setup as in Experiment 1, we report the results of the L2 loss through time for Active-NeuralUCB.

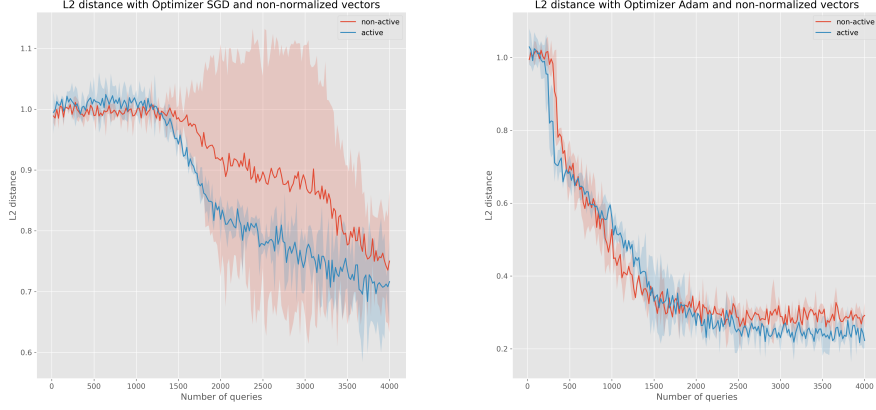


Figure 3: Performance SVM on non-normalized learned features Adam vs SGD

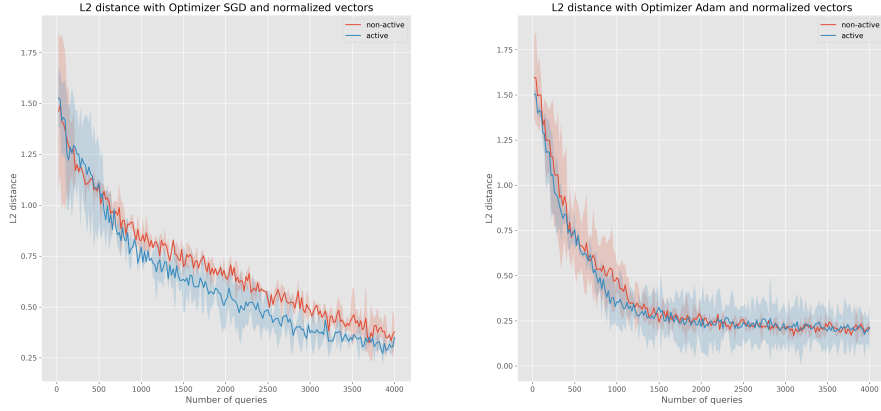


Figure 4: Performance SVM on normalized learned features Adam vs SGD

7.3 Experiment 3: Regret of OnSim-NeuralUCB

We observe the performance of OnSim-NeuralUCB by comparing the regret of the algorithm with the regret of an epsilon-greedy neural network with the same architecture and hyperparameters. We report the results only Adam with normalized and unnormalized versions. Like in the previous experiment all graphs contain the git mean and a 95% confidence interval. We also compare various values for the exploration parameter.²

8 Conclusion

References

- [1] Wei Chu, Lihong Li, Lev Reyzin, and Robert Schapire. Contextual bandits with linear payoff functions. In Geoffrey Gordon, David Dunson, and Miroslav Dudík, editors, *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, volume 15 of

²The appendix contains a more thorough graph with even more values.

Proceedings of Machine Learning Research, pages 208–214, Fort Lauderdale, FL, USA, 11–13 Apr 2011. PMLR.

- [2] Arthur Jacot, Franck Gabriel, and Clément Hongler. Neural tangent kernel: Convergence and generalization in neural networks. 2018.
- [3] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization, 2014.
- [4] Dongruo Zhou, Lihong Li, and Quanquan Gu. Neural contextual bandits with ucb-based exploration, 2019.

A Appendix

Algorithm 5: Train

Input : ϵ learning rate, E number of epochs, dataset $\{(x_i^{k_i}, y_i^{k_i}, r_i)\}_{i=1}^t$, b_s batch size

```
1 optimizer  $\leftarrow$  SGD( $\epsilon, \theta$ );
2 for  $j = 1, \dots, E$  do
3    $\mathcal{D} \leftarrow \{(x_1^{k_1}, y_1^{k_1}, r_1), \dots, (x_t^{k_t}, y_t^{k_t}, r_t)\}$ ;
4   while  $\mathcal{D}$  is not empty do
5     Sample minibatch  $M$  of size  $b_s$  from the training examples;
6      $l = \frac{1}{M} \sum_{(x_i^{k_i}, y_i^{k_i}, r_i) \in M} (\phi(x_i^{k_i}, y_i^{k_i}) - r_i)^2$ ;
7      $\theta \leftarrow$  update using optimizer with loss  $l$ ;
8     Remove  $M$  from  $\mathcal{D}$ ;
9   end
10 end
```
